

Notes for cs452

Kernel Description

Bill Cowan
University of Waterloo

I. Programmer's Model

I.1. Task Creation

This section comprises the part of the API used to support the first part of kernel development. Create is non-trivial; the remainder do little more than providing support for testing.

I.1. a. Create

Create – instantiate a task.

Synopsis.

```
int Create( int priority, void (*code) ( ) )
```

Description.

Create allocates and initializes a task descriptor, using the given priority, and the given function pointer as a pointer to the entry point of executable code, essentially a function with no arguments and no return value. When Create returns the task descriptor has all the state needed to run the task, the task's stack has been suitably initialized, and the task has been entered into its ready queue so that it will run the next time it is scheduled.

Returns.

`tid` the positive integer task id of the newly created task. The task id must be unique, in the sense that no task has, will have or has had the same task id.

-1 if the priority is invalid.

-2 if the kernel is out of task descriptors.

Comment.

You might want to do some rough tests to ensure that the function pointer argument is valid: it's not obvious to me which of the many possible tests would be most useful.

I.1. b. MyTid

MyTid – return my task id.

Synopsis

```
int MyTid( )
```

Description

MyTid returns the task id of the calling task.

Returns

`tid` the positive integer task id of the task that calls it.

Comments.

Errors should be impossible!

I.1. c. MyParentTid

Name

MyParentTid – return the task id of the task that created the calling task.

Synopsis

```
int MyParentTid( )
```

Description

MyParentTid returns the task id of the task that created the calling task. This will be problematic only if the task has exited or been destroyed, in which case the return value is implementation-dependent.

Returns

`tid` the task id of the task that created the calling task.

Comments

The return value is implementation-dependent if the parent has exited, has been destroyed, or is in the process of being destroyed.

I.1. d. Pass

Name

Pass – cease execution, remaining ready to run.

Synopsis

```
void Pass( )
```

Description

Pass causes a task to stop executing. The task is moved to the end of its priority queue, and will resume executing when next scheduled.

Comments

Like every entry to the kernel, Pass reschedules.

I.1. e. Exit

Name

Exit – terminate execution forever.

Synopsis

```
void Exit( )
```

Description

Exit causes a task to cease execution permanently. It is removed from all priority queues, send queues, receive queues and awaitEvent queues. Resources owned by the task, primarily its memory and task descriptor are not reclaimed.

Returns

Exit does not return. If a point occurs where all tasks have exited the kernel should return cleanly to RedBoot.

1.1. f. Destroy

Please see the separate document for Destroy. Re-using resources is complicated.

1.2. Inter-task Communication

This subsection comprises the primitives required for message passing, which are needed for the second part of the kernel.

1.2. a. Send

Name

Send – send a message to a specific task and obtain the corresponding response.

Synopsis

```
int Send(int tid, char *msg, int msglen, char *reply, int rplen)
```

Description

Send sends a message to another task and receives a reply. The message, in a buffer in the sending task's address space is copied to the address space of the task to which it is sent by the kernel. Send supplies a buffer into which the reply is to be copied, and the size of the buffer so that the kernel can detect overflow. When Send returns without error it is guaranteed that the message has been received, and that a reply has been sent, not necessarily by the same task. If either the message or the reply is a string it is necessary that the length should include the terminating null.

The kernel will not overflow the reply buffer. The caller is expected to compare the return value to the size of the reply buffer. If part of the reply is missing the return value will exceed the size of the supplied reply buffer.

There is no guarantee that Send will return. If, for example, the task to which the message is directed never calls Receive, Send never returns and the sending task remains blocked forever.

Send has a passing resemblance, and no more, to remote procedure call.

Returns

- >-1 The size of the message responded by the replying task. The message is less than or equal to the size of the buffer provided for it. Longer responses are truncated.
- 1 The reply message was truncated.
- 2 The task id supplied is not the task id of an existing task.
- 3 The send-receive-reply transaction could not be completed.

1.2. b. Receive

Name

Receive - receive a message from any task.

Synopsis

```
int Receive( int *tid, char *msg, int msglen )
```

Description

Receive blocks until a message is sent to the caller, then returns with the message in its message buffer and tid set to the task id of the task that sent the message. Messages sent before Receive is called are retained in a send queue, from which they are received in first-come, first-served order.

The argument msg must point to a buffer at least as large as msglen. If the size of the message received exceeds msglen, no overflow occurs and the buffer will contain the first msglen characters of the message sent.

The caller is expected to compare the return value, which contains the size of the message that was sent, to determine whether or not the message is complete, and to act accordingly.

Returns

- >-1 The size of the message received, which is less than or equal to the size of the message buffer supplied. Longer messages are truncated.
- 1 The message is truncated.

1.2. c. Reply

Name

Reply - reply to a message.

Synopsis

```
int Reply( int tid, char *reply, int rplen )
```

Description

Reply sends a reply to a task that previously sent a message. When it returns without error, the reply has been copied into the sender's address space. The calling task and the sender return at the same logical time, so whichever is of higher priority runs first. If they are of the same priority the sender runs first.

Returns

- 0 The reply succeeded.
- 1 The message was truncated.
- 2 The task id is not the task id of an existing task.
- 3 The task id is not the task id of a reply blocked task.

1.3. Name Server

This section comprises the API of the name server.

1.3. a. RegisterAs

Name

RegisterAs – register a name with the name server.

Synopsis

```
int RegisterAs( char *name )
```

Description

RegisterAs registers the task id of the caller under the given name.

On return without error it is guaranteed that all WhoIs calls by any task will return the task id of the caller until the registration is overwritten.

If another task has already registered with the given name its registration is overwritten.

A single task may register under several different names, but each name is assigned to a single task.

RegisterAs is actually a wrapper covering a send to the name server.

Returns

- 0 Success.
- 1 The nameserver task id inside the wrapper is invalid.

1.3. b. WhoIs

Name

WhoIs - query the nameserver.

Synopsis

```
int WhoIs( char *name )
```

Description

WhoIs asks the nameserver for the task id of the task that is registered under the given name.

Whether WhoIs blocks waiting for a registration or returns with an error if no task is registered under the given name is implementation-dependent.

There is guaranteed to be a unique task id associated with each registered name, but the registered task may change at any time after a call to WhoIs.

WhoIs is actually a wrapper covering a send to the nameserver.

Returns

- tid The task id of the registered task.
- 1 The nameserver task id inside the wrapper is invalid.

1.4. Interrupt Processing

One primitive is required for generic interrupt processing.

1.4. a. AwaitEvent

Name

AwaitEvent - wait for an external event.

Synopsis

```
int AwaitEvent( int eventid )
```

Description

AwaitEvent blocks until the event identified by eventid occurs then returns, with volatile data, if any.

Returns

- >-1 volatile data, in the form of a positive integer.
- 1 invalid event.
- 2 corrupted volatile data.

I.5. Clock Server

I.5. a. Delay

Name

Delay - wait for a given amount of time.

Synopsis

```
int Delay( int tid, int ticks )
```

Description

Delay returns after the given number of ticks has elapsed. How long after is not guaranteed because the caller may have to wait on higher priority tasks.

Delay is (almost) identical to Pass if ticks is zero or negative.

The size of a tick is normally application dependent. In cs452 this term it is 10 milliseconds.

Delay is actually a wrapper for a send to the clock server.

Returns

- 0 Success.
- 1 The clock server task id is invalid.
- 2 The delay was zero or negative.

I.5. b. Time

Name

Time - give the time since clock server start up.

Synopsis

```
int Time( int tid )
```

Description

Time returns the number of ticks since the clock server was created and initialized.

With a 10 millisecond tick and a 32-bit unsigned int for the time wraparound is almost a million hours, plenty of time for your demo.

Time is actually a wrapper for a send to the clock server. The argument is the tid of the clock server.

Returns

- >-1 The time in ticks since the clock server initialized.
- 1 The clock server task id is invalid.

I.5. c. DelayUntil

Name

DelayUntil - wait until a time.

Synopsis

```
int DelayUntil( int tid, int ticks )
```

Description

Delay returns when the time since clock server initialization is greater than the given number of ticks. How long after is not guaranteed because the caller may have to wait on higher priority tasks.

DelayUntil(tid, Time(tid) + ticks) may differ from Delay(tid, ticks) by a small amount.

The size of a tick is normally application dependent. In cs452 this term it is 10 milliseconds, the time in which a train at top speed travels about 5 millimetres.

DelayUntil is actually a wrapper for a send to the clock server.

Returns

- 0 Success.
- 1 The clock server task id is invalid.
- 2 The delay was zero or negative.

1.6. Input/Output

1.6. a. Getc

Name

Getc - get a character from a UART.

Synopsis

```
int Getc( int tid, int uart )
```

Description

Getc returns next unreturned character from the given UART. The first argument is the task id of the appropriate server.

How communication errors are handled is implementation-dependent.

Getc is actually a wrapper for a send to the appropriate server.

Returns

- >0 Success.
- 1 The server task id is not the task id of an existing task.

1.6. b. Putc

Name

Putc - transmit a character from the given UART.

Synopsis

```
int Putc( int tid, int uart, char ch )
```

Description

Putc queues the given character for transmission by the given UART. On return the only guarantee is that the character has been queued. Whether it has been transmitted or received is not guaranteed.

How configuration errors are handled is implementation-dependent.

Putc is actually a wrapper for a send to the serial server.

Returns

- 0 Success.
- 1 The server task id is not the task id of an existing task.

II. Algorithms and Data Structures

All algorithms must be constant-time, in the sense that you can bound them above at a reasonable level. ‘Reasonable’ means on a time scale appropriate for a train application.

All data must be either static or on the stack. There is no heap; there is no dynamic memory allocation, at least as you know it in C, C++ or Java running under Unix.

II.1. Task Descriptors

II.1. a. Basics

The most important data structure in the kernel is the array of task descriptors (TDs), which is allocated on the kernel stack during kernel initialization. Every existing task has a TD allocated to it. A TD normally includes at least

1. a task identifier (tid), which is unique among all tasks, past, present and future,
2. a pointer to the TD of the task that created it, its parent,
3. the task’s priority,
4. a pointer to the TD of the next task in the task’s ready queue,
5. a point to the TD of the next task on the task’s send queue,
6. the task’s current run state, and
7. the task’s current stack pointer.

The remainder of the task’s state is saved on the task’s stack. A few items are sometimes in the TD, sometimes on the stack. They include

- the task’s return value, and
- the task’s SPSR.

II.1. b. Comments

1. When Destroy is not implemented the task id can be its array index because TDs are not re-used. When Destroy is implemented a better task id model is needed.
2. A task is in one of the following run states.
 - i. Active. The task that has just run, is running, or is about to run. Scheduling, which happens near the end of kernel processing, changes the active task. On a single processor only one task can be active at a time.
 - ii. Ready. The task is ready to be activated.
 - iii. Zombie. The task will never again run, but still retains its resources: memory, TD, etc.
 - iv. Send-blocked. The task has executed Receive, and is waiting for a task to sent to it.
 - v. Receive-blocked. The task has executed Send, and is waiting for the message to be received.
 - vi. Reply-blocked. The task has executed Send and its message has been received, but it has not received a reply.
 - vii. Event-blocked. The task has executed AwaitEvent, but the event on which it is waiting has not occurred.

The first three of these states are needed for task creation; the next three are needed for message passing; and the seventh is needed for hardware interrupts.

3. The parent is the active task when a task is being created. This entails that the variable storing the active task is written only by the scheduler.
4. A task has memory reserved for it when it is created, which it uses for its stack. The values of its registers are placed on the stack when it is not executing.
5. Each task has, when it is running, a program status register (CPSR), which is saved when it is interrupted and re-installed when the task is next activated.
6. Tasks usually enter the kernel with requests for service. Many tasks must be provided with return values that indicate the result of the request. Because a task may not be rescheduled immediately the return value must be saved.

II.2. Scheduling

II.2. a. Priorities

Scheduling is done using static priorities: higher priority tasks that are ready to execute are guaranteed to execute before lower priority ones.

1. The number of priorities is implementation-dependent.
2. There can be more than one task at any priority. As tasks become ready to run, they are put on the end of their ready queue. The next task to run is always the highest priority task that is ready. If more than one task is ready at the highest priority, then the one at the head of the ready queue runs.
3. A task instance may not be at more than one priority.

II.2. b. Round-Robin

While a task can have only a single priority; several tasks may have the same priority. They are scheduled round-robin. Each time scheduling occurs at a priority, the least recently readied task is scheduled, requiring something like a first-in, first-out queue.

II.2. c. No Ready Tasks

Occasionally, or even frequently, all ready queues are empty with one or more tasks Event-Blocked. In this case there are two things you might do.

1. You can have an idle task at the lowest priority, which does anything you want while it waits to be interrupted.
2. You can put the CPU into the halt state.

The second consumes less power, but power consumption is unimportant in cs452.

II.2. d. Exiting from the Kernel

During the first two parts of kernel development the kernel should exit cleanly to RedBoot when there are no tasks in the ready queues. Once interrupts are introduced the kernel should exit cleanly to RedBoot when there are no Event-blocked tasks and no tasks in the ready queues. Kernels do not normally exit.

II.3. Context Switching

Those context switches into the kernel that occur when a running task requests a kernel service must be implemented using the ARM SWI instruction. Context switches generated by interrupts are defined by the architecture. The ARM architecture offers a variety of methods for exiting the kernel into a task. Choose whichever suits you.

II.4. Interrupts

The kernel should run with interrupts disabled. Thus, as little as possible should be done by the kernel because the time of the slowest kernel primitive must be added to the worst case response time of every action.

AwaitEvent requires a table of event ids, which is, in essence a catalogue of hardware events to which the kernel and servers are able to provide a response. This is an awkward intrusion of hardware configuration into an otherwise clean operating system abstraction. There are many different ways of handling this issue within operating system design. Arguments between proponents of different approaches are, for the most part, religious, which means that although the arguments are heated and often personal, they only rarely change the minds of the arguers. Fortunately for you, the hardware environment in which you work is circumscribed and well-defined, so that you can create a small registry of events, provide a reasonably consistent set of responses, and not think too much about the general problem. But we will discuss the general problem in class, just because it is quite pervasive, affecting many different areas of computer science.

If you happen to have a really good solution, and can persuade the world to take it up you are well on your way to your first billion.

II.5. Message Formats

In most kernels and applications few, if any, messages are character strings; almost all messages contain the contents of a struct. Type checking might be handy, but can only be performed at run time because there is no restriction on which task can send to which other task. You can get pretty close to run-time type checking by giving every structure a field, the value of which is its type. What you do when you discover a mismatch is the hard part. Can you recover at run-time? Or do you need to stop and reprogram.

You will, no doubt, find that early design of a set of message structures, combined with discipline in your task structure, is important in defending you from bugs that can take a long time to find.